

Fast License Checker - Build Plan

Executive Summary

This document provides a complete, step-by-step action plan for building a blazing-fast, Rust-based CLI tool and library for checking and fixing license headers in source files. The implementation follows the **Modular Monolith** architecture pattern with strict type safety, comprehensive testing, and CI/CD quality gates.

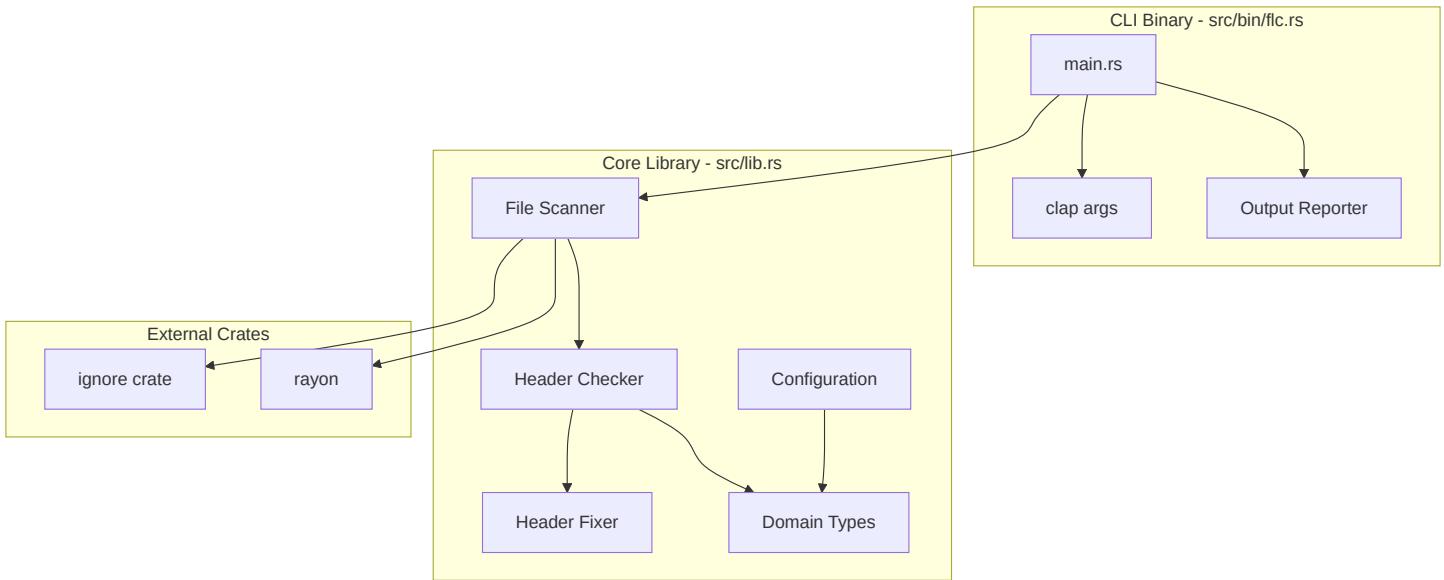
Target Performance: 100,000 files scanned in under 1 second (warm cache)

Table of Contents

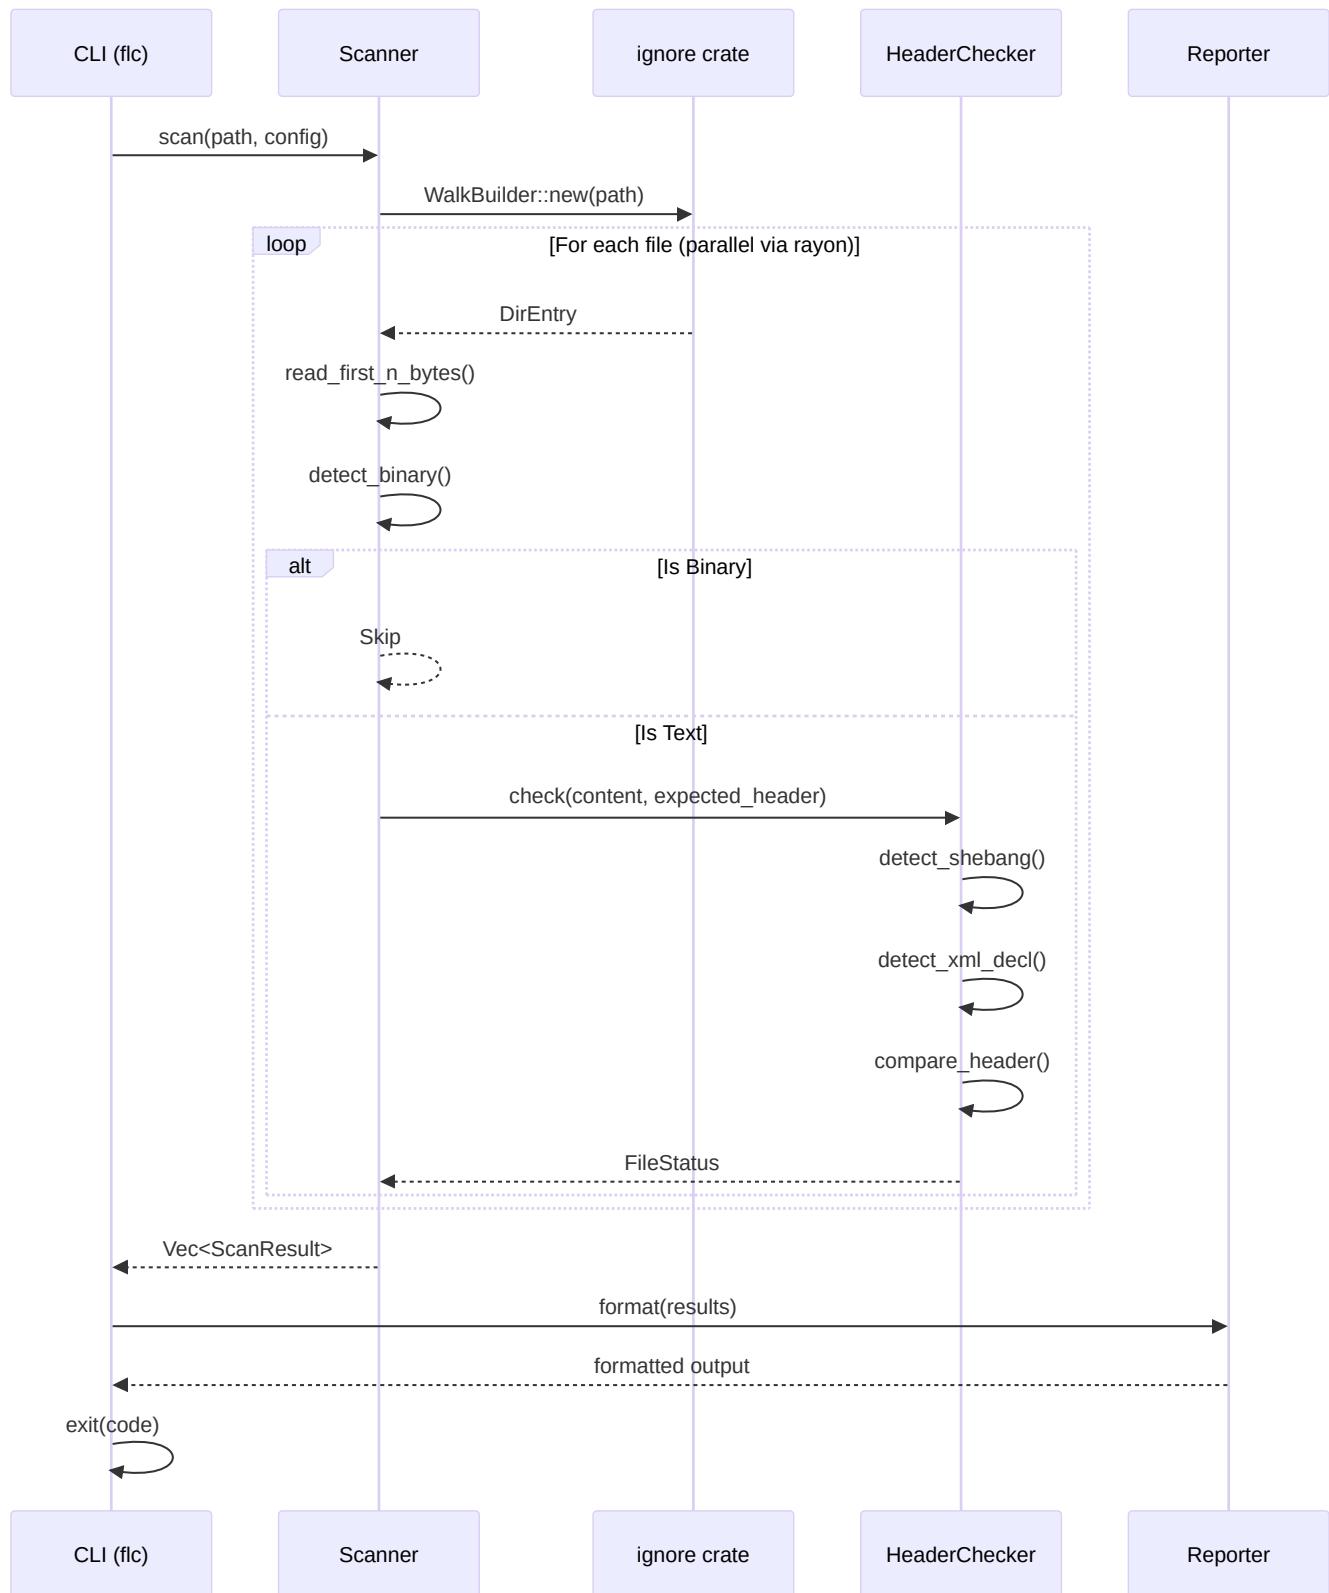
1. [Architecture Overview](#)
2. [Project Structure](#)
3. [Phase 1: Bootstrap](#)
4. [Phase 2: RFC and Design](#)
5. [Phase 3: Core Library Implementation](#)
6. [Phase 4: CLI Implementation](#)
7. [Phase 5: Testing](#)
8. [Phase 6: CI/CD Setup](#)
9. [Cursor Workflow Summary](#)
10. [Key Architecture Decisions](#)
11. [Validation Checkpoints](#)

Architecture Overview

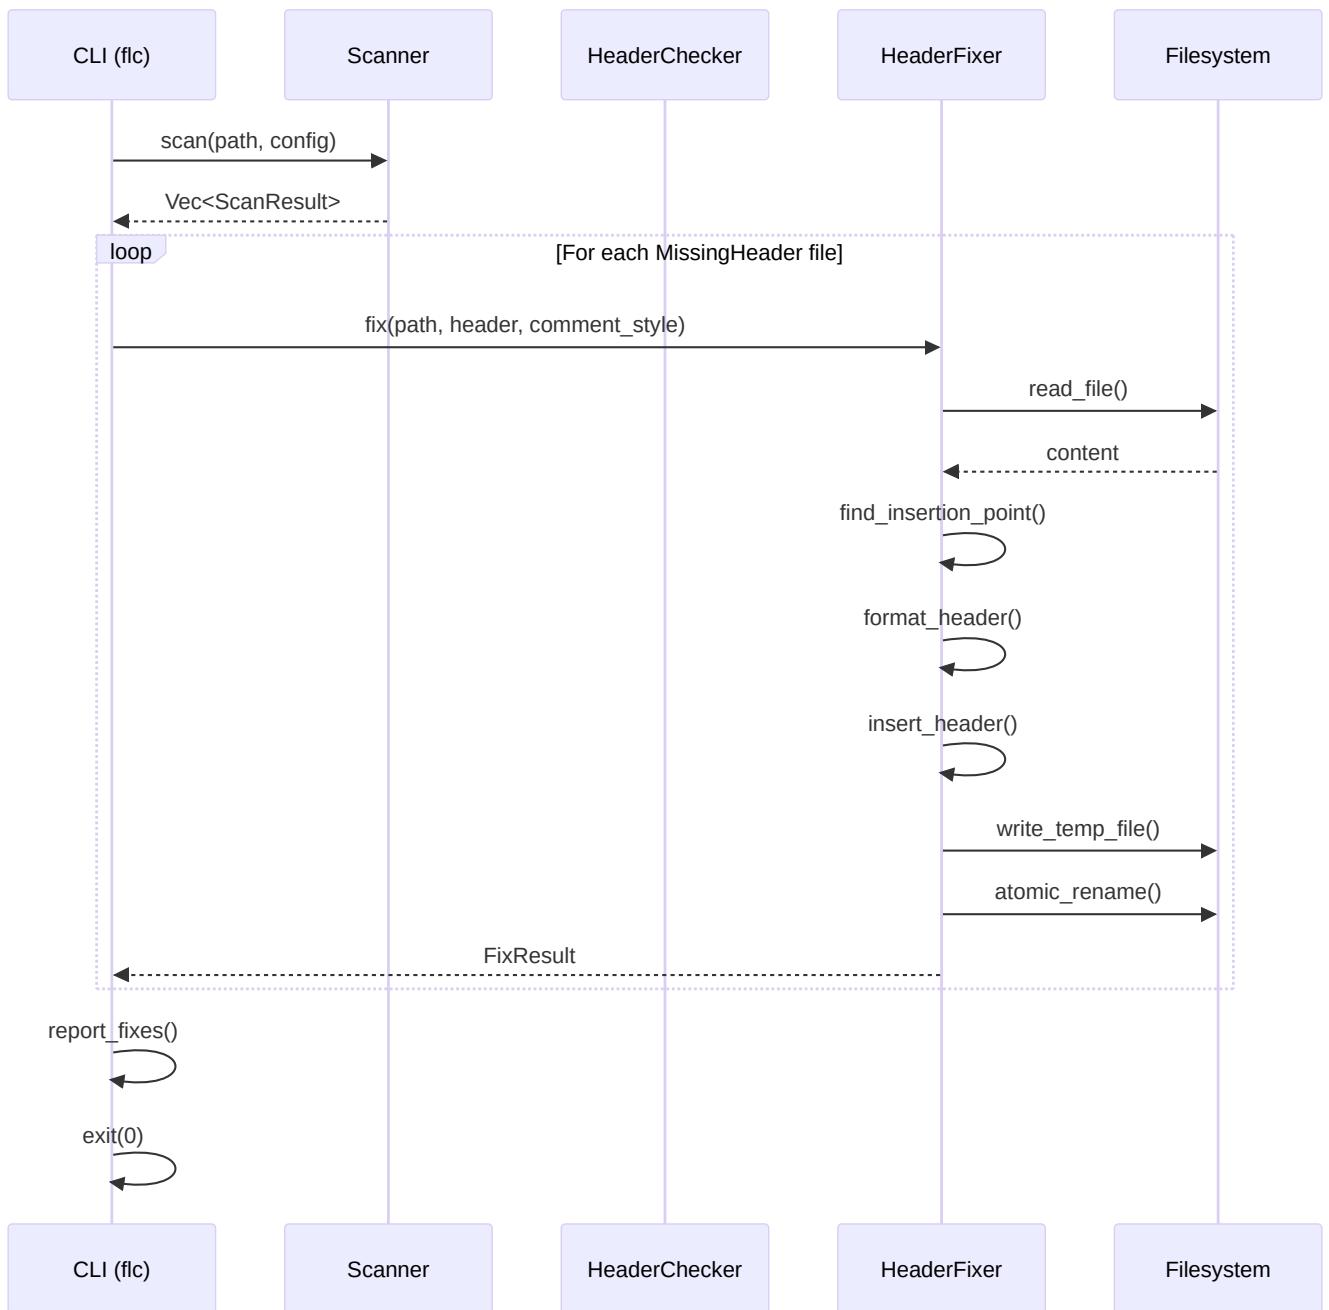
This project follows the **Modular Monolith** pattern, adapted for a CLI tool:



Data Flow - Scan Mode



Data Flow - Fix Mode



Project Structure

```
fast-license-checker/
├── .cursorrules                         # AI coding guidelines
├── .github/
│   ├── workflows/
│   │   └── ci.yml                         # Quality gate CI
│   └── scripts/
│       └── check_primitives.sh           # Ban raw primitives in core
├── Cargo.toml                            # Package manifest
├── Cargo.lock
├── clippy.toml                           # Strict linting rules
├── deny.toml                             # Supply chain security policy
├── rustfmt.toml                          # Formatting rules
└── README.md

|
├── docs/
│   ├── inputs/                            # Business requirements
│   │   └── fast-license-checker-requirements.txt
│   ├── rfcs/                             # Technical design docs
│   │   ├── 0000-build-plan.md            # This document
│   │   └── 0001-license-checker.md      # Feature RFC
│   └── adr/                              # Architecture decisions
        └── 001-use-ignore-crate.md

|
└── src/
    ├── lib.rs                            # Library entry point
    ├── design_sketch.rs                 # Type skeleton (temporary)
    |
    ├── config/                           # Configuration handling
    │   ├── mod.rs
    │   ├── types.rs                      # Config structs
    │   └── loader.rs                     # TOML/CLI loading
    |
    ├── scanner/                          # File walking & filtering
    │   ├── mod.rs
    │   ├── walker.rs                     # ignore crate wrapper
    │   └── filter.rs                     # Binary detection
    |
    ├── checker/                          # Header detection & validation
    │   ├── mod.rs
    │   └── detector.rs                  # Header presence check
```

```

|   |   └── validator.rs          # Similarity matching
|   └── prelude.rs               # Shebang/XML detection
|
|   └── fixer/                  # Header insertion logic
|       ├── mod.rs
|       ├── inserter.rs          # Header formatting
|       └── writer.rs            # Atomic file writes
|
|   └── types/                  # Domain NewTypes
|       ├── mod.rs
|       ├── file_types.rs        # FilePath, FileExtension
|       ├── header_types.rs      # LicenseHeader, CommentStyle
|       └── results.rs           # FileStatus, ScanResult
|
|   └── error.rs                # Typed errors (thiserror)
|
└── src/bin/
    └── flc.rs                  # CLI entry point (clap + anyhow)
|
└── tests/
    ├── common/
    |   └── mod.rs               # Test fixtures & helpers
    └── integration/
        ├── scan_mode.rs          # Scan mode e2e tests
        ├── fix_mode.rs            # Fix mode e2e tests
        └── edge_cases.rs          # Shebang, binary, etc.
|
└── benches/
    └── scan_performance.rs     # 100k file benchmark
|
└── snapshots/
    └── *.snap                  # insta test snapshots

```

Phase 1: Bootstrap (Day 1)

Step 1.1: Initialize Project

Cursor Prompt:

Create a new Rust project for a CLI tool called "fast-license-checker" with the following requirements:

1. Cargo.toml with:

- Package name: fast-license-checker
- Binary name: flc
- Edition 2021
- License: MIT OR Apache-2.0
- Workspace lints for clippy (deny unwrap_used, expect_used, panic)

2. Dependencies:

- ignore = "0.4" (from ripgrep, for file walking)
- rayon = "1.10" (parallel processing)
- clap = { version = "4", features = ["derive", "env"] }
- thiserror = "2"
- anyhow = "1"
- tracing = "0.1"
- tracing-subscriber = { version = "0.3", features = ["fmt", "env-filter"] }
- serde = { version = "1", features = ["derive"] }
- toml = "0.8" (for config files)
- memchr = "2" (fast byte searching)

3. Dev dependencies:

- tempfile = "3" (for test fixtures)
- criterion = { version = "0.5", features = ["html_reports"] }
- proptest = "1"
- insta = { version = "1", features = ["json"] }

4. Create empty src/lib.rs and src/bin/flc.rs with minimal structure

5. Add [[bench]] section for criterion:

```
[[bench]]
name = "scan_performance"
harness = false
```

Expected Output: Valid Cargo.toml, empty lib.rs and bin/flc.rs that compile

Step 1.2: Configuration Files

Cursor Prompt:

Create the following configuration files for the fast-license-checker project:

1. clippy.toml with strict rules enforcing reliability:

```
[lints.clippy]
# Tenet 6: Reliability - Ban crashes
unwrap_used = "deny"
expect_used = "deny"
panic = "deny"
indexing_slicing = "deny"

# Tenet 2: Correctness - Prevent silent bugs
arithmetic_side_effects = "warn"

# Tenet 7: Simplicity - Clean code
wildcard_imports = "deny"
enum_glob_use = "deny"
```

2. deny.toml for supply chain security:

```
[advisories]
vulnerability = "deny"
unmaintained = "warn"

[licenses]
unlicensed = "deny"
allow = ["MIT", "Apache-2.0", "BSD-2-Clause", "BSD-3-Clause", "ISC", "MPL-2.0", "Unic
copyleft = "deny"

[bans]
multiple-versions = "warn"
```

3. rustfmt.toml with consistent formatting:

```
max_width = 100
use_small_heuristics = "Max"
edition = "2021"
newline_style = "Unix"
use_field_init_shorthand = true
use_try_shorthand = true
```

Validation: Run cargo fmt --check , cargo clippy , cargo deny check

Step 1.3: Create .cursorrules

Cursor Prompt:

Create .cursorrules file for fast-license-checker project encoding these principles:

```
# Project: Fast License Checker - Blazing Fast License Header Verification
# Role: You are a senior Rust engineer building a high-performance CLI tool
```

Core Principles (7 Tenets)

1. **Security (Defense in Depth)**

- Parse, don't validate: wrap all inputs in validated NewTypes
- Never trust file content - detect binary, handle encoding errors
- Use secrecy crate if handling sensitive data

2. **Performance (Zero-Cost Abstractions)**

- All NewTypes use #[repr(transparent)]
- Use byte slices (&[u8]) over String where possible
- Stream files, never load entire repo into RAM
- Stop reading after max_header_bytes (default 8KB)
- Use memchr for fast byte searching

3. **Scalability (Parallel by Default)**

- Use rayon for parallel file processing
- Use ignore crate's parallel walker
- Design for 100k+ files

4. **Readability (Bus Factor = 1)**

- Every public fn/struct must have /// doc comments
- Explain WHY, not just WHAT
- No magic; explicit > clever

5. **Observability (The Eyes)**

- Use #[tracing::instrument] on all public functions
- Use tracing macros (debug!, info!, warn!, error!), NOT println!
- Include relevant context in spans

6. **Reliability (Resilience)**

- ZERO .unwrap() or .expect() in production code
- All errors typed with thiserror in library
- Use anyhow only in CLI binary for user-friendly errors
- Never panic—return Result instead

7. **Simplicity (Maintainability)**

- Prefer stable, boring crates: ignore, rayon, clap, serde
- Single binary, library + CLI architecture

- Let clippy and rustfmt decide style disputes

Architecture Constraints

```
### Core Library (src/lib.rs and submodules)
- **Allowed imports**: serde, thiserror, tracing, rayon, ignore, memchr
- **FORBIDDEN**: clap, anyhow (CLI-only)
- Purpose: Pure logic for scanning, checking, fixing
- Testable in isolation
```

CLI Binary (src/bin/flc.rs)

```
- **Allowed imports**: clap, anyhow, tracing-subscriber, core library
- Purpose: Argument parsing, output formatting, exit codes
- Thin wrapper around library
```

Code Style & Patterns

NewTypes (Zero-Cost Wrappers)

Always wrap domain concepts:

```
```rust
#[repr(transparent)]
#[derive(Debug, Clone, PartialEq, Eq, Hash)]
pub struct LicenseHeader(String);

impl LicenseHeader {
 pub fn new(s: impl Into<String>) -> Result<Self, ValidationError> {
 let s = s.into();
 let trimmed = s.trim();
 if trimmed.is_empty() {
 return Err(ValidationError::EmptyHeader);
 }
 Ok(Self(trimmed.to_owned()))
 }

 pub fn as_str(&self) -> &str { &self.0 }
 pub fn as_bytes(&self) -> &[u8] { self.0.as_bytes() }
}
```

## Error Handling

Use thiserror for library errors:

```

#[derive(Debug, thiserror::Error)]
pub enum CheckerError {
 #[error("IO error reading {path}: {source}")]
 Io {
 path: std::path::PathBuf,
 #[source] source: std::io::Error
 },
 #[error("Binary file detected: {0}")]
 BinaryFile(std::path::PathBuf),
}

```

## Instrumentation

Add tracing to functions:

```

#[tracing::instrument(skip(content), fields(path = %path.display()))]
pub fn check_header(
 path: &Path,
 content: &[u8],
 expected: &LicenseHeader,
) -> Result<FileStatus, CheckerError> {
 tracing::debug!(content_len = content.len(), "Checking header");
 // ...
}

```

## Forbidden Patterns (CI will fail)

- ✗ .unwrap() in production code
- ✗ .expect() in production code
- ✗ panic!() macros
- ✗ Array indexing arr[i] (use .get() )
- ✗ use crate::\*;

(explicit imports only)
- ✗ println!() / eprintln!() - use tracing macros
- ✗ Raw primitives in public function signatures

# File Handling Safety Rules

1. **Binary Detection:** Check first 8KB for NULL bytes before processing
2. **Size Limits:** Never read more than max\_header\_bytes for checks
3. **Encoding:** Assume UTF-8, skip files that fail validation
4. **Shebangs:** Detect `#!` at start, insert header AFTER
5. **XML:** Detect `<?xml` at start, insert header AFTER
6. **Atomic Writes:** Write to temp file, then rename
7. **Idempotency:** Never duplicate headers - check before fixing

## Testing Requirements

- Every public function has at least one unit test
- Use proptest for invariant testing
- Use insta for snapshot testing of output formats
- Integration tests with tempfile for real filesystem
- Benchmark with criterion for performance regression

**\*\*Validation:\*\*** File exists and is comprehensive

---

**### Step 1.4: Create .vscode/settings.json {#step-14-create-vscode-settingsjson }**

**\*\*Cursor Prompt:\*\***

Create .vscode/settings.json for Cursor IDE with rust-analyzer configuration:

```
{
 "rust-analyzer.check.command": "clippy",
 "rust-analyzer.check.extraArgs": ["--", "-D", "warnings"],
 "rust-analyzer.cargo.features": "all",
 "rust-analyzer.procMacro.enable": true,
 "rust-analyzer.diagnostics.experimental.enable": true,
 "editor.formatOnSave": true,
 "editor.rulers": [100],
 "[rust)": {
```

```
"editor.defaultFormatter": "rust-lang.rust-analyzer"
}
}

Phase 2: RFC and Design (Day 2) {#phase-2-rfc-and-design-day-2 }

Step 2.1: Generate RFC Document {#step-21-generate-rfc-document }

Cursor Prompt:
```

@docs/inputs/fast-license-checker-requirements.txt

Create RFC at docs/rfcs/0001-license-checker.md following this structure:

# RFC 0001: Fast License Checker

## Executive Summary

### Problem Statement:

Manual license header verification doesn't scale. Existing tools are slow, poorly configured, or don't respect .gitignore.

### Proposed Solution:

A blazing-fast CLI tool that:

- Scans directories respecting .gitignore
- Checks for license headers with fuzzy matching
- Auto-fixes missing headers (with shebang/XML awareness)
- Provides multiple output formats for CI integration

### Goals:

- Scan 100,000 files in < 1 second (warm cache)
- Zero false positives on binary files
- Idempotent fix mode
- Exit code 0 = clean, 1 = issues found

## Non-Goals (v1):

- License compatibility checking (GPL vs MIT)
- Fetching licenses from internet
- OCR/image scanning

# User Stories

## Happy Path: Clean Repository

As a maintainer, I want to run `f1c .` and see "All 1,234 files have valid license headers" with exit code 0.

## Sad Path 1: Missing Headers

As a DevOps engineer, when files are missing headers, I want a clear list of paths and exit code 1 so CI blocks the PR.

## Sad Path 2: Shebang Files

As a maintainer, when I run fix mode on a shell script with `#!/bin/bash`, the license header should be inserted on line 2, preserving the shebang.

## Sad Path 3: Binary Files

As a user, when scanning a directory with images/binaries, those files should be automatically skipped without errors.

## Sad Path 4: Malformed Headers

As a compliance officer, when a file has a header that's 80% similar but not exact, I want a warning to manually review it.

## Sad Path 5: Large Files

As a user, when scanning a repo with large log files, the tool should only read the first N bytes and not hang.

# System Design

## Sequence Diagram - Scan Mode

[Include mermaid diagram from architecture overview]

## Sequence Diagram - Fix Mode

[Include mermaid diagram from architecture overview]

## Component Architecture

[Include mermaid diagram showing modules]

## Architectural Decision Records

Decision	Context	Trade-offs
Use <code>ignore crate</code>	Need <code>.gitignore</code> support, parallel walking	Adds dependency, but battle-tested in ripgrep
Use <code>rayon</code> over <code>tokio</code>	File scanning is CPU-bound, not IO-bound	Can't use async, but simpler mental model
Use byte slices <code>&amp;[u8]</code>	Binary detection, performance	Must handle encoding explicitly
Use <code>thiserror + anyhow</code>	Typed errors in lib, friendly in CLI	Two error handling styles
Atomic writes via rename	Prevent corruption on crash	Requires temp file creation
Max 8KB header check	Don't read entire large files	Could miss headers placed later (rare)

# Data Models (Skeleton)

```
// NewTypes
pub struct FilePath(PathBuf);
pub struct LicenseHeader(String);
pub struct CommentStyle { prefix: String, suffix: Option<String> }
pub struct FileExtension(String);
pub struct MaxHeaderBytes(usize);

// Enums
pub enum FileStatus {
 HasHeader,
 MissingHeader,
 MalformedHeader { similarity: u8 },
 Skipped { reason: SkipReason },
}

pub enum SkipReason {
 Binary,
 Empty,
 Gitignored,
 TooLarge,
 UnsupportedEncoding,
}

pub enum ScanMode {
 Check,
 Fix,
}

// Result types
pub struct ScanResult {
 pub path: FilePath,
 pub status: FileStatus,
}

pub struct FixResult {
 pub path: FilePath,
 pub action: FixAction,
}

pub enum FixAction {
```

```
 Fixed,
 AlreadyHasHeader,
 Skipped { reason: SkipReason },
 Failed { error: String },
}

}
```

## Implementation Phases

1. **Phase 1:** Bootstrap - project structure, configs, .cursorrules
2. **Phase 2:** Design - RFC, design\_sketch.rs that compiles
3. **Phase 3:** Core types - NewTypes, errors, config
4. **Phase 4:** Scanner - file walking with ignore crate
5. **Phase 5:** Checker - header detection with fuzzy matching
6. **Phase 6:** Fixer - header insertion with shebang awareness
7. **Phase 7:** CLI - clap interface, output formats
8. **Phase 8:** Testing - unit, integration, property-based
9. **Phase 9:** Benchmarks - criterion, 100k file target
10. **Phase 10:** CI/CD - GitHub Actions, release builds

---

```
Step 2.2: Create Design Sketch {#step-22-create-design-sketch }
```

\*\*Cursor Prompt:\*\*

@docs/rfcs/0001-license-checker.md

Create design sketch at src/design\_sketch.rs that compiles with cargo check .

This file validates our type design before full implementation. Include:

1. All NewTypes with #[repr(transparent)] where appropriate:
  - FilePath(PathBuf) - wraps PathBuf
  - LicenseHeader(String) - non-empty validated
  - CommentStyle { prefix: String, suffix: Option }
  - FileExtension(String) - lowercase, no dot
  - MaxHeaderBytes(usize) - minimum 256
  - SimilarityScore(u8) - 0-100

## 2. All enums:

- FileStatus { HasHeader, MissingHeader, MalformedHeader { similarity: SimilarityScore }, Skipped(SkipReason) }
- SkipReason { Binary, Empty, Gitignored, TooLarge, UnsupportedEncoding, NoCommentStyle }
- ScanMode { Check, Fix }
- FixAction { Fixed, AlreadyHasHeader, Skipped(SkipReason), Failed(String) }

## 3. Result structs:

- ScanResult { path: FilePath, status: FileStatus }
- FixResult { path: FilePath, action: FixAction }
- ScanSummary { total: usize, passed: usize, failed: usize, skipped: usize, duration: Duration }

## 4. Trait definitions:

- trait HeaderChecker { fn check(&self, content: &[u8], expected: &LicenseHeader, style: &CommentStyle) -> FileStatus; }
- trait HeaderFixer { fn fix(&self, content: &[u8], header: &LicenseHeader, style: &CommentStyle) -> Result<Vec, FixerError>; }

## 5. Error types skeleton:

- enum LicenseCheckerError (IO, Config, Validation variants)
- enum ConfigError (NotFound, Parse, MissingField variants)
- enum FixerError (BinaryFile, WriteError, IdempotencyViolation variants)

Add doc comments to all public items. Run `cargo check` to validate.

\*\*Validation:\*\* `cargo check` passes with `design_sketch.rs`

---

## Phase 3: Core Library Implementation (Days 3-5) {#phase-3-core-library-implementation}

### Step 3.1: Implement Domain Types {#step-31-implement-domain-types }

\*\*Cursor Prompt:\*\*

@src/design\_sketch.rs

Implement the NewTypes in `src/types/mod.rs`, `src/types/file_types.rs`, `src/types/header_types.rs`, and `src/types/results.rs`.

Move types from `design_sketch.rs` to their proper modules:

1. src/types/file\_types.rs:

- FilePath with constructor that validates path exists (optional, for fix mode)
- FileExtension with lowercase normalization, no leading dot
- MaxHeaderBytes with minimum of 256, default of 8192

2. src/types/header\_types.rs:

- LicenseHeader with non-empty validation after trim
- CommentStyle with common presets (rust\_style(), python\_style(), html\_style(), etc.)
- SimilarityScore clamped to 0-100

3. src/types/results.rs:

- FileStatus enum
- SkipReason enum
- ScanResult struct
- FixResult struct
- FixAction enum
- ScanSummary struct

Requirements:

- Add #[repr(transparent)] to single-field wrappers
- Derive: Debug, Clone, PartialEq, Eq where appropriate
- Add Serialize, Deserialize for JSON output support
- Implement Display for error messages
- Add AsRef and From where ergonomic
- Comprehensive /// doc comments

4. src/types/mod.rs:

- Re-export all public types
- Module documentation explaining the type system

\*\*Validation:\*\* `cargo check`, `cargo doc --open` to verify docs

---

### Step 3.2: Implement Error Types {#step-32-implement-error-types }

\*\*Cursor Prompt:\*\*

Implement typed errors in src/error.rs using thiserror:

```
use std::path::PathBuf;
use thiserror::Error;

/// Top-level error type for the license checker library
#[derive(Debug, Error)]
pub enum LicenseCheckerError {
 #[error("Configuration error: {0}")]
 Config(#[from] ConfigError),
 #[error("Scanner error: {0}")]
 Scanner(#[from] ScannerError),
 #[error("Checker error: {0}")]
 Checker(#[from] CheckerError),
 #[error("Fixer error: {0}")]
 Fixer(#[from] FixerError),
}

/// Configuration-related errors
#[derive(Debug, Error)]
pub enum ConfigError {
 #[error("Config file not found: {0}")]
 NotFound(PathBuf),
 #[error("Failed to parse config: {0}")]
 Parse(#[from] toml::de::Error),
 #[error("Missing required field: {field}")]
 MissingField { field: &'static str },
 #[error("Invalid value for {field}: {message}")]
 InvalidValue { field: &'static str, message: String },
}

/// File scanning errors
#[derive(Debug, Error)]
pub enum ScannerError {
 #[error("Failed to walk directory {path}: {source}")]
 WalkError {
 path: PathBuf,
 #[source] source: ignore::Error
 },
}
```

```
#[error("IO error reading {path}: {source}")]
Io {
 path: PathBuf,
 #[source] source: std::io::Error
},
}

/// Header checking errors
#[derive(Debug, Error)]
pub enum CheckerError {
 #[error("IO error reading {path}: {source}")]
 Io {
 path: PathBuf,
 #[source] source: std::io::Error
 },
 #[error("File appears to be binary: {0}")]
 BinaryFile(PathBuf),
 #[error("Unsupported encoding in file: {0}")]
 UnsupportedEncoding(PathBuf),
}

/// Header fixing errors
#[derive(Debug, Error)]
pub enum FixerError {
 #[error("Cannot fix binary file: {0}")]
 BinaryFile(PathBuf),
 #[error("Failed to write {path}: {source}")]
 WriteError {
 path: PathBuf,
 #[source] source: std::io::Error
 },
 #[error("Header already exists in {0} - refusing to duplicate")]
 IdempotencyViolation(PathBuf),
 #[error("Malformed header detected in {path} (similarity: {similarity}%) - manual re MalformedHeader { path: PathBuf, similarity: u8 },
}
```

```
/// Validation errors for NewTypes
#[derive(Debug, Error)]
pub enum ValidationError {
 #[error("License header cannot be empty")]
 EmptyHeader,
 #[error("File extension cannot be empty")]
 EmptyExtension,
 #[error("MaxHeaderBytes must be at least 256, got {0}")]
 HeaderBytesTooSmall(usize),
 #[error("Similarity score must be 0-100, got {0}")]
 InvalidSimilarity(u8),
}
```

Add impl blocks for converting to anyhow::Error in the CLI.

---

```
Step 3.3: Implement Configuration Module {#step-33-implement-configuration-module }
```

\*\*Cursor Prompt:\*\*

Implement configuration in src/config/mod.rs, src/config/types.rs, and src/config/loader.rs:

1. src/config/types.rs - Config struct:

```
use std::collections::HashMap;
use std::path::PathBuf;
use serde::{Deserialize, Serialize};

/// Main configuration for the license checker
#[derive(Debug, Clone, Serialize, Deserialize)]
#[serde(default)]
pub struct Config {
 /// The license header text to check for
 pub license_header: String,

 /// Comment styles per file extension
 pub comment_styles: HashMap<String, CommentStyleConfig>,

 /// Additional glob patterns to ignore (beyond .gitignore)
 pub ignore_patterns: Vec<String>,

 /// Maximum bytes to read from file start for header check
 pub max_header_bytes: usize,

 /// Skip empty files (0 bytes)
 pub skip_empty_files: bool,

 /// Number of parallel jobs (None = num_cpus)
 pub parallel_jobs: Option<usize>,

 /// Similarity threshold for malformed header detection (0-100)
 pub similarity_threshold: u8,
}

#[derive(Debug, Clone, Serialize, Deserialize)]
pub struct CommentStyleConfig {
 pub prefix: String,
 #[serde(default)]
 pub suffix: Option<String>,
}

impl Default for Config {
 fn default() -> Self {
 Self {
 license_header: String::new(),
 comment_styles: default_comment_styles(),
 ignore_patterns: vec![],
```

```
 max_header_bytes: 8192,
 skip_empty_files: true,
 parallel_jobs: None,
 similarity_threshold: 70,
 }
}

fn default_comment_styles() -> HashMap<String, CommentStyleConfig> {
 // Include mappings for:
 // - .rs, .js, .ts, .c, .cpp, .h, .java, .go, .swift, .kt, .scala -> // prefix
 // - .py, .rb, .sh, .bash, .zsh, .yaml, .yml, .toml, .pl, .pm -> # prefix
 // - .html, .xml, .svg, .vue -> <!-- prefix, --> suffix
 // - .css, .scss, .less -> /* prefix, */ suffix
 // - .sql -> -- prefix
 // - .lua -> -- prefix
 // - .hs -> -- prefix
 // - .erl -> % prefix
 // - .ex, .exs -> # prefix
 // - .clj -> ;; prefix
 // - .lisp -> ;; prefix
 // - .vim -> " prefix
 // - .bat, .cmd -> REM prefix
}
```

2. src/config/loader.rs - Loading logic:

```

use std::path::Path;

/// Load config from file, CLI args, and environment
#[tracing::instrument]
pub fn load_config(
 config_path: Option<&Path>,
 cli_overrides: CliOverrides,
) -> Result<Config, ConfigError> {
 // 1. Start with defaults
 // 2. Load from file if exists
 // 3. Apply CLI overrides
 // 4. Apply environment variables (FLC_HEADER, FLC_MAX_BYTES)
 // 5. Validate final config
}

pub struct CliOverrides {
 pub license_header: Option<String>,
 pub license_file: Option<PathBuf>,
 pub parallel_jobs: Option<usize>,
 pub max_header_bytes: Option<usize>,
}

```

### 3. src/config/mod.rs - Re-exports and module docs

---

### Step 3.4: Implement File Scanner {#step-34-implement-file-scanner }

**\*\*Cursor Prompt:\*\***

Implement file scanner in src/scanner/mod.rs, src/scanner/walker.rs, and src/scanner/filter.rs:

#### 1. src/scanner/walker.rs - File walking with ignore crate:

```

use ignore::WalkBuilder;
use rayon::prelude::*;
use std::path::Path;

/// File walker that respects .gitignore and provides parallel iteration
pub struct FileWalker {
 root: PathBuf,
 additional_ignores: Vec<String>,
 parallel_jobs: usize,
}

impl FileWalker {
 pub fn new(root: impl AsRef<Path>) -> Self { ... }

 pub fn with_ignores(mut self, patterns: Vec<String>) -> Self { ... }

 pub fn with_parallelism(mut self, jobs: usize) -> Self { ... }

 /// Walk all files, yielding DirEntry for each
 #[tracing::instrument(skip(self))]
 pub fn walk(&self) -> impl ParallelIterator<Item = Result<WalkEntry, ScannerError>>
 WalkBuilder::new(&self.root)
 .hidden(true) // Skip hidden files
 .git_ignore(true) // Respect .gitignore
 .git_global(true) // Respect global gitignore
 .git_exclude(true) // Respect .git/info/exclude
 .threads(self.parallel_jobs)
 .build_parallel()
 // Convert to parallel iterator via rayon
 }
}

pub struct WalkEntry {
 pub path: PathBuf,
 pub depth: usize,
 pub file_type: FileType,
}

```

2. src/scanner/filter.rs - Binary detection:

```

/// Detect if content is likely binary (contains NULL bytes)
#[tracing::instrument(skip(content))]
pub fn is_binary(content: &[u8]) -> bool {
 // Use memchr for fast NULL byte search
 memchr::memchr(0, content).is_some()
}

/// Detect if content is valid UTF-8
pub fn is_valid_utf8(content: &[u8]) -> bool {
 std::str::from_utf8(content).is_ok()
}

/// Check if file should be skipped based on content
pub fn should_skip(content: &[u8], config: &Config) -> Option<SkipReason> {
 if content.is_empty() && config.skip_empty_files {
 return Some(SkipReason::Empty);
 }
 if is_binary(content) {
 return Some(SkipReason::Binary);
 }
 if !is_valid_utf8(content) {
 return Some(SkipReason::UnsupportedEncoding);
 }
 None
}

```

3. src/scanner/mod.rs - Main scanner interface:

```

/// Main scanner that coordinates walking and checking
pub struct Scanner {
 walker: FileWalker,
 checker: HeaderChecker,
 config: Config,
}

impl Scanner {
 #[tracing::instrument(skip(config))]
 pub fn new(root: impl AsRef<Path>, config: Config) -> Result<Self, ScannerError> {

 /// Scan all files and return results
 #[tracing::instrument(skip(self))]
 pub fn scan(&self) -> ScanSummary {
 let start = Instant::now();

 let results: Vec<ScanResult> = self.walker.walk()
 .filter_map(|entry| entry.ok())
 .filter(|entry| entry.file_type.is_file())
 .par_bridge()
 .map(|entry| self.check_file(&entry))
 .collect();

 ScanSummary::from_results(results, start.elapsed())
 }

 fn check_file(&self, entry: &WalkEntry) -> ScanResult { ... }
 }
}

```

---

### Step 3.5: Implement Header Checker {#step-35-implement-header-checker }

\*\*Cursor Prompt:\*\*

Implement header checker in src/checker/mod.rs, src/checker/detector.rs, src/checker/validator.rs, and src/checker/prelude.rs:

1. src/checker/prelude.rs - Shebang and XML detection:

```

/// Detect shebang line and return byte offset after it
pub fn detect_shebang(content: &[u8]) -> Option<usize> {
 if content.starts_with(b"#!") {
 // Find end of first line
 memchr::memchr(b'\n', content).map(|pos| pos + 1)
 } else {
 None
 }
}

/// Detect XML declaration and return byte offset after it
pub fn detect_xml_declaration(content: &[u8]) -> Option<usize> {
 if content.starts_with(b"<?xml") {
 // Find closing ?>
 content.windows(2)
 .position(|w| w == b"?>")
 .map(|pos| {
 // Skip past ?> and any following newline
 let end = pos + 2;
 if content.get(end) == Some(&b'\n') {
 end + 1
 } else {
 end
 }
 })
 } else {
 None
 }
}

/// Get the byte offset where header should start (after shebang/xml)
pub fn header_start_offset(content: &[u8]) -> usize {
 detect_shebang(content)
 .or_else(|| detect_xml_declaration(content))
 .unwrap_or(0)
}

```

2. src/checker/detector.rs - Header presence detection:

```

/// Check if the expected header is present in content
#[tracing::instrument(skip(content, expected))]
pub fn detect_header(
 content: &[u8],
 expected: &LicenseHeader,
 style: &CommentStyle,
) -> HeaderMatch {
 let start_offset = header_start_offset(content);
 let search_region = &content[start_offset..];

 // Format expected header with comment style
 let formatted_header = format_header_for_search(expected, style);

 // Exact match
 if search_region.starts_with(formatted_header.as_bytes()) {
 return HeaderMatch::Exact;
 }

 // Fuzzy match for malformed headers
 if let Some(similarity) = fuzzy_match(search_region, &formatted_header) {
 if similarity >= 70 {
 return HeaderMatch::Fuzzy { similarity };
 }
 }
}

HeaderMatch::None
}

pub enum HeaderMatch {
 Exact,
 Fuzzy { similarity: u8 },
 None,
}

```

3. src/checker/validator.rs - Similarity calculation:

```

/// Calculate similarity between two byte slices (0-100)
pub fn calculate_similarity(a: &[u8], b: &[u8]) -> u8 {
 // Simple approach: compare first N bytes of each
 // Use Levenshtein distance ratio for more accuracy

 let max_len = a.len().max(b.len());
 if max_len == 0 {
 return 100;
 }

 let matching = a.iter()
 .zip(b.iter())
 .filter(|(x, y)| x == y)
 .count();

 ((matching * 100) / max_len) as u8
}

```

#### 4. src/checker/mod.rs - Main checker interface:

```

pub struct HeaderChecker {
 expected_header: LicenseHeader,
 comment_styles: HashMap<FileExtension, CommentStyle>,
 max_bytes: MaxHeaderBytes,
}

impl HeaderChecker {
 #[tracing::instrument(skip(self, path))]
 pub fn check_file(&self, path: &Path) -> Result<FileStatus, CheckerError> {
 // 1. Read first max_bytes of file
 // 2. Check for binary/encoding issues
 // 3. Get appropriate comment style for extension
 // 4. Detect header presence
 // 5. Return FileStatus
 }
}

```

- - -

```
Step 3.6: Implement Header Fixer {#step-36-implement-header-fixer }
```

**\*\*Cursor Prompt:\*\***

Implement header fixer in src/fixer/mod.rs, src/fixer/inserter.rs, and src/fixer/writer.rs:

1. src/fixer/inserter.rs - Header formatting and insertion:

```
/// Format license header with appropriate comment style
pub fn format_header(
 header: &LicenseHeader,
 style: &CommentStyle,
) -> String {
 let mut result = String::new();

 for line in header.as_str().lines() {
 result.push_str(&style.prefix);
 if !line.is_empty() {
 result.push(' ');
 result.push_str(line);
 }
 if let Some(ref suffix) = style.suffix {
 result.push(' ');
 result.push_str(suffix);
 }
 result.push('\n');
 }

 // Add blank line after header
 result.push('\n');

 result
}

/// Insert header into content at the correct position
pub fn insert_header(
 content: &[u8],
 header: &LicenseHeader,
 style: &CommentStyle,
) -> Result<Vec<u8>, FixerError> {
 let insert_offset = header_start_offset(content);
 let formatted = format_header(header, style);

 let mut result = Vec::with_capacity(content.len() + formatted.len());

 // Copy content before insertion point (shebang/xml)
 result.extend_from_slice(&content[..insert_offset]);

 // Insert header
 result.extend_from_slice(formatted.as_bytes());
}
```

```
// Copy rest of content
result.extend_from_slice(&content[insert_offset..]);

Ok(result)
}
```

2. src/fixer/writer.rs - Atomic file writing:

```
use std::fs::{self, File};
use std::io::Write;
use std::path::Path;

/// Write content to file atomically (write temp, then rename)
#[tracing::instrument(skip(content))]
pub fn write_atomic(path: &Path, content: &[u8]) -> Result<(), FixerError> {
 let parent = path.parent().ok_or_else(|| FixerError::WriteError {
 path: path.to_path_buf(),
 source: std::io::Error::new(
 std::io::ErrorKind::InvalidInput,
 "Path has no parent directory"
),
 })?;
 // Create temp file in same directory (for atomic rename)
 let temp_path = parent.join(format!("{}.{}.tmp",
 path.file_name()
 .and_then(|n| n.to_str())
 .unwrap_or("file")
));
 // Write to temp file
 {
 let mut file = File::create(&temp_path).map_err(|e| FixerError::WriteError {
 path: temp_path.clone(),
 source: e,
 })?;
 file.write_all(content).map_err(|e| FixerError::WriteError {
 path: temp_path.clone(),
 source: e,
 })?;
 file.sync_all().map_err(|e| FixerError::WriteError {
 path: temp_path.clone(),
 source: e,
 })?;
 }
 // Atomic rename
 fs::rename(&temp_path, path).map_err(|e| FixerError::WriteError {
 path: path.to_path_buf(),
 })
}
```

```
 source: e,
})?;

tracing::info!(path = %path.display(), "Fixed file");

ok()
}
```

### 3. src/fixer/mod.rs - Main fixer interface:

```
pub struct HeaderFixer {
 header: LicenseHeader,
 comment_styles: HashMap<FileExtension, CommentStyle>,
 checker: HeaderChecker,
}

impl HeaderFixer {
 /// Fix a single file, adding header if missing
 #[tracing::instrument(skip(self))]
 pub fn fix_file(&self, path: &Path) -> Result<FixResult, FixerError> {
 // 1. Read entire file
 let content = fs::read(path).map_err(|e| FixerError::WriteError {
 path: path.to_path_buf(),
 source: e,
 })?;

 // 2. Check if binary
 if is_binary(&content) {
 return Ok(FixResult {
 path: FilePath::new(path),
 action: FixAction::Skipped(SkipReason::Binary),
 });
 }

 // 3. Check current status
 let status = self.checker.check_content(&content, path)?;

 match status {
 FileStatus::HasHeader => {
 Ok(FixResult {
 path: FilePath::new(path),
 action: FixAction::AlreadyHasHeader,
 })
 }
 FileStatus::MalformedHeader { similarity } => {
 // Don't auto-fix malformed headers
 Err(FixerError::MalformedHeader {
 path: path.to_path_buf(),
 similarity,
 })
 }
 FileStatus::MissingHeader => {
 // Get comment style for this file
 }
 }
 }
}
```

```
 let ext = get_extension(path);
 let style = self.comment_styles
 .get(&ext)
 .ok_or_else(|| FixerError::WriteError {
 path: path.to_path_buf(),
 source: std::io::Error::new(
 std::io::ErrorKind::Unsupported,
 format!("No comment style for extension: {}", ext)
),
 })?;

 // Insert header
 let new_content = insert_header(&content, &self.header, style)?;

 // Write atomically
 write_atomic(path, &new_content)?;

 Ok(FixResult {
 path: FilePath::new(path),
 action: FixAction::Fixed,
 })
}

FileStatus::Skipped(reason) => {
 Ok(FixResult {
 path: FilePath::new(path),
 action: FixAction::Skipped(reason),
 })
}
}
```

## Phase 4: CLI Implementation (Day 6) {#phase-4-cli-implementation-day-6}

### Step 4.1: Implement CLI Arguments {#step-41-implement-cli-arguments}

**\*\*Cursor Prompt:\*\***

Implement CLI in src/bin/flc.rs using clap derive::

```
use clap::{Parser, ValueEnum};
use std::path::PathBuf;

/// Fast License Checker - Blazing fast license header verification
#[derive(Parser, Debug)]
#[command(name = "flc")]
#[command(author, version, about, long_about = None)]
pub struct Cli {
 /// Directory or file to scan
 #[arg(default_value = ".")]
 pub path: PathBuf,

 /// Run in fix mode (add missing headers)
 #[arg(short, long)]
 pub fix: bool,

 /// Path to file containing license header text
 #[arg(short = 'l', long = "license")]
 pub license_file: Option<PathBuf>,

 /// License header text (alternative to --license)
 #[arg(long = "header", conflicts_with = "license_file")]
 pub header_text: Option<String>,

 /// Config file path
 #[arg(short, long, default_value = ".license-checker.toml")]
 pub config: PathBuf,

 /// Number of parallel jobs (default: number of CPUS)
 #[arg(short, long, env = "FLC_JOBS")]
 pub jobs: Option<usize>,

 /// Maximum bytes to read for header check
 #[arg(long, default_value = "8192", env = "FLC_MAX_BYTES")]
 pub max_bytes: usize,

 /// Output format
 #[arg(short, long, value_enum, default_value = "text")]
 pub output: OutputFormat,

 /// Increase verbosity (-v, -vv, -vvv)
 #[arg(short, long, action = clap::ArgAction::Count)]
 pub verbose: u8,
```

```
/// Suppress all output except errors
#[arg(short, long, conflicts_with = "verbose")]
pub quiet: bool,

/// Don't use colors in output
#[arg(long)]
pub no_color: bool,
}

#[derive(Debug, Clone, Copy, ValueEnum)]
pub enum OutputFormat {
 /// Human-readable text output
 Text,
 /// JSON output for programmatic use
 Json,
 /// GitHub Actions annotation format
 Github,
}
```

---

```
Step 4.2: Implement Main Entry Point {#step-42-implement-main-entry-point }
```

\*\*Cursor Prompt:\*\*

Implement main function in src/bin/flc.rs:

```
use anyhow::Context, Result;
use clap::Parser;
use tracing_subscriber::{EnvFilter, fmt};

mod cli;
mod output;

use cli::{Cli, OutputFormat};
use fast_license_checker::{Config, Scanner, Fixer, ScanSummary};

fn main() -> Result<()> {
 let cli = Cli::parse();

 // Initialize tracing based on verbosity
 init_tracing(cli.verbose, cli.quiet)?;

 // Load configuration
 let config = load_config(&cli)
 .context("Failed to load configuration")?;

 tracing::debug!(?config, "Loaded configuration");

 // Validate license header is provided
 if config.license_header.is_empty() {
 anyhow::bail!(
 "No license header provided. Use --license <file> or --header <text>, \
 or add 'license_header' to your config file."
);
 }

 // Run scan or fix
 let result = if cli.fix {
 run_fix_mode(&cli, &config)
 } else {
 run_scan_mode(&cli, &config)
 };

 // Handle result and exit
 match result {
 Ok(summary) => {
 output::print_summary(&summary, cli.output, !cli.no_color);

 if summary.failed > 0 {

```

```
 std::process::exit(1);
 }
 Ok(())
}
Err(e) => {
 tracing::error!(error = %e, "Operation failed");
 Err(e)
}
}

fn init_tracing(verbose: u8, quiet: bool) -> Result<()> {
 let level = if quiet {
 "error"
 } else {
 match verbose {
 0 => "info",
 1 => "debug",
 _ => "trace",
 }
 };
 let filter = EnvFilter::try_from_default_env()
 .unwrap_or_else(|_| EnvFilter::new(level));

 fmt()
 .with_env_filter(filter)
 .with_target(false)
 .init();

 Ok(())
}

fn run_scan_mode(cli: &Cli, config: &Config) -> Result<ScanSummary> {
 let scanner = Scanner::new(&cli.path, config.clone())
 .context("Failed to create scanner")?;

 let summary = scanner.scan();

 Ok(summary)
}

fn run_fix_mode(cli: &Cli, config: &Config) -> Result<ScanSummary> {
```

```
let fixer = Fixer::new(&cli.path, config.clone())
 .context("Failed to create fixer")?;

let summary = fixer.fix_all()
 .context("Fix operation failed")?;

Ok(summary)
}

fn load_config(cli: &Cli) -> Result<Config> {
 // Implementation using config module
}

Step 4.3: Implement Output Formatters {#step-43-implement-output-formatters }

Cursor Prompt:
```

Create src/bin/output.rs with output formatting:

```
use fast_license_checker::{ScanSummary, ScanResult, FileStatus};
use crate::cli::OutputFormat;

pub fn print_summary(summary: &ScanSummary, format: OutputFormat, color: bool) {
 match format {
 OutputFormat::Text => print_text(summary, color),
 OutputFormat::Json => print_json(summary),
 OutputFormat::Github => print_github(summary),
 }
}

fn print_text(summary: &ScanSummary, color: bool) {
 use std::io::{self, Write};
 let mut stdout = io::stdout().lock();

 // Header
 if color {
 writeln!(stdout, "\n\x1b[1mFast License Checker Results\x1b[0m").ok();
 } else {
 writeln!(stdout, "\nFast License Checker Results").ok();
 }
 writeln!(stdout, "{}", "=" .repeat(40)).ok();

 // Stats
 writeln!(stdout, "Total files: {}", summary.total).ok();
 writeln!(stdout, "With header: {}", summary.passed).ok();

 if color && summary.failed > 0 {
 writeln!(stdout, "Missing header: \x1b[31m{}\x1b[0m", summary.failed).ok();
 } else {
 writeln!(stdout, "Missing header: {}", summary.failed).ok();
 }

 writeln!(stdout, "Skipped: {}", summary.skipped).ok();
 writeln!(stdout, "Duration: {:.2}s", summary.duration.as_secs_f64()).ok();
 writeln!(stdout, "Throughput: {:.0} files/sec",
 summary.total as f64 / summary.duration.as_secs_f64())
 .ok();

 // List failures
 if !summary.failures.is_empty() {
 writeln!(stdout, "\nFiles missing license header:").ok();
 for result in &summary.failures {
```

```

 if color {
 writeln!(stdout, " \x1b[31mx\x1b[0m {}", result.path).ok();
 } else {
 writeln!(stdout, " - {}", result.path).ok();
 }
 }

// List malformed
if !summary.malformed.is_empty() {
 writeln!(stdout, "\nFiles with malformed headers (manual review required)").ok();
 for result in &summary.malformed {
 if let FileStatus::MalformedHeader { similarity } = result.status {
 if color {
 writeln!(stdout, " \x1b[33m\x1b[0m {} ({}% similar)",
 result.path, similarity).ok();
 } else {
 writeln!(stdout, " - {} ({}% similar)", result.path, similarity).ok();
 }
 }
 }
}

// Final status
writeln!(stdout).ok();
if summary.failed == 0 && summary.malformed.is_empty() {
 if color {
 writeln!(stdout, "\x1b[32m✓ All files have valid license headers\x1b[0m").ok();
 } else {
 writeln!(stdout, "OK: All files have valid license headers").ok();
 }
} else {
 if color {
 writeln!(stdout, "\x1b[31mx Some files need attention\x1b[0m").ok();
 } else {
 writeln!(stdout, "FAIL: Some files need attention").ok();
 }
}
}

fn print_json(summary: &ScanSummary) {
 // Serialize to JSON using serde_json
 let json = serde_json::json!({

```

```

 "total": summary.total,
 "passed": summary.passed,
 "failed": summary.failed,
 "skipped": summary.skipped,
 "duration_secs": summary.duration.as_secs_f64(),
 "failures": summary.failures.iter().map(|r| &r.path).collect::<Vec<_>>(),
 "malformed": summary.malformed.iter().map(|r| {
 serde_json::json!({
 "path": &r.path,
 "similarity": match r.status {
 FileStatus::MalformedHeader { similarity } => similarity,
 _ => 0
 }
 })
 }).collect::<Vec<_>>(),
});

println!("{}", serde_json::to_string_pretty(&json).unwrap_or_default());
}

fn print_github(summary: &ScanSummary) {
 // GitHub Actions annotation format
 for result in &summary.failures {
 println!("::error file={}::Missing license header", result.path);
 }

 for result in &summary.malformed {
 if let FileStatus::MalformedHeader { similarity } = result.status {
 println!("::warning file={}::Malformed license header ({}% similar)",
 result.path, similarity);
 }
 }
}

// Summary annotation
if summary.failed > 0 {
 println!("::error::License check failed: {} files missing headers", summary.fai-
}
}

```

- - -

```
Phase 5: Testing (Day 7) {#phase-5-testing-day-7 }

Step 5.1: Create Unit Tests {#step-51-create-unit-tests }

Cursor Prompt:
```

Create unit tests for all modules. For each module, create a tests submodule or [tests.rs](#) file.

1. src/types/tests.rs:

```
#[cfg(test)]
mod tests {
 use super::*;

 use proptest::prelude::*;

 #[test]
 fn license_header_rejects_empty() {
 assert!(LicenseHeader::new("").is_err());
 assert!(LicenseHeader::new(" ").is_err());
 assert!(LicenseHeader::new("\n\t").is_err());
 }

 #[test]
 fn license_header_trims_whitespace() {
 let header = LicenseHeader::new(" MIT License ").unwrap();
 assert_eq!(header.as_str(), "MIT License");
 }

 #[test]
 fn file_extension_normalizes_lowercase() {
 let ext = FileExtension::new("RS").unwrap();
 assert_eq!(ext.as_str(), "rs");
 }

 #[test]
 fn file_extension_removes_dot() {
 let ext = FileExtension::new(".py").unwrap();
 assert_eq!(ext.as_str(), "py");
 }

 #[test]
 fn max_header_bytes_minimum() {
 assert!(MaxHeaderBytes::new(100).is_err());
 assert!(MaxHeaderBytes::new(256).is_ok());
 }

 proptest! {
 #[test]
 fn license_header_never_panics(s in ".*") {
 let _ = LicenseHeader::new(s);
 }

 #[test]
 }
}
```

```
fn similarity_score_clamped(n in 0u8..=255u8) {
 let score = SimilarityScore::new(n);
 assert!(score.is_ok() == (n <= 100));
}
}
```

2. src/checker/tests.rs:

```
#[cfg(test)]
mod tests {
 use super::*;

 const MIT_HEADER: &str = "Copyright (c) 2024\nMIT License";

 fn rust_style() -> CommentStyle {
 CommentStyle { prefix: "/".into(), suffix: None }
 }

 #[test]
 fn detects_header_at_start() {
 let content = b"// Copyright (c) 2024\n// MIT License\nfn main() {}";
 let header = LicenseHeader::new(MIT_HEADER).unwrap();

 let status = detect_header(content, &header, &rust_style());
 assert!(matches!(status, HeaderMatch::Exact));
 }

 #[test]
 fn detects_header_after_shebang() {
 let content = b"#!/bin/bash\n# Copyright (c) 2024\n# MIT License\nnecho hello";
 let header = LicenseHeader::new(MIT_HEADER).unwrap();
 let style = CommentStyle { prefix: "#".into(), suffix: None };

 let status = detect_header(content, &header, &style);
 assert!(matches!(status, HeaderMatch::Exact));
 }

 #[test]
 fn detects_header_after_xml_declaration() {
 let content = b"<?xml version=\"1.0\"?>\n<!-- Copyright (c) 2024 -->\n<!-- MIT !";
 let header = LicenseHeader::new(MIT_HEADER).unwrap();
 let style = CommentStyle { prefix: "<!--".into(), suffix: Some("-->".into()) };

 let status = detect_header(content, &header, &style);
 assert!(matches!(status, HeaderMatch::Exact));
 }

 #[test]
 fn returns_missing_when_no_header() {
 let content = b"fn main() { println!("Hello"); }";
 let header = LicenseHeader::new(MIT_HEADER).unwrap();
 }
}
```

```
 let status = detect_header(content, &header, &rust_style());
 assert!(matches!(status, HeaderMatch::None));
}

#[test]
fn handles_crlf_line_endings() {
 let content = b"// Copyright (c) 2024\r\n// MIT License\r\n\r\nfn main() {}";
 let header = LicenseHeader::new(MIT_HEADER).unwrap();

 let status = detect_header(content, &header, &rust_style());
 assert!(matches!(status, HeaderMatch::Exact));
}
}
```

3. src/fixer/tests.rs:

```

#[cfg(test)]
mod tests {
 use super::*;

#[test]
fn inserts_header_at_start() {
 let content = b"fn main() {}";
 let header = LicenseHeader::new("MIT").unwrap();
 let style = CommentStyle { prefix: "//".into(), suffix: None };

 let result = insert_header(content, &header, &style).unwrap();

 assert!(result.starts_with(b"// MIT\n"));
 assert!(result.ends_with(b"fn main() {}"));
}

#[test]
fn preserves_shebang() {
 let content = b"#!/bin/bash\necho hello";
 let header = LicenseHeader::new("MIT").unwrap();
 let style = CommentStyle { prefix: "#".into(), suffix: None };

 let result = insert_header(content, &header, &style).unwrap();

 assert!(result.starts_with(b"#!/bin/bash\n"));
 assert!(result.contains(&b"# MIT"[..]));
}

#[test]
fn idempotency_check_blocks_duplicate() {
 // First insertion
 let content = b"fn main() {}";
 let header = LicenseHeader::new("MIT").unwrap();
 let style = CommentStyle { prefix: "//".into(), suffix: None };

 let fixed = insert_header(content, &header, &style).unwrap();

 // Checker should now detect header
 let status = detect_header(&fixed, &header, &style);
 assert!(matches!(status, HeaderMatch::Exact));
}

}

```

---

```
Step 5.2: Create Integration Tests {#step-52-create-integration-tests }
```

\*\*Cursor Prompt:\*\*

Create integration tests in tests/integration/:

1. tests/common/mod.rs - Test fixtures:

```
use std::fs;
use std::path::Path;
use tempfile::TempDir;

pub struct TestFixture {
 pub dir: TempDir,
}

impl TestFixture {
 pub fn new() -> Self {
 Self { dir: TempDir::new().unwrap() }
 }

 pub fn create_file(&self, name: &str, content: &str) -> std::path::PathBuf {
 let path = self.dir.path().join(name);
 if let Some(parent) = path.parent() {
 fs::create_dir_all(parent).unwrap();
 }
 fs::write(&path, content).unwrap();
 path
 }

 pub fn create_gitignore(&self, patterns: &[&str]) {
 let content = patterns.join("\n");
 self.create_file(".gitignore", &content);
 }

 pub fn path(&self) -> &Path {
 self.dir.path()
 }
}
```

2. tests/integration/scan\_mode.rs:

```
mod common;

use std::process::Command;
use common::TestFixture;

#[test]
fn scan_empty_directory_returns_zero() {
 let fixture = TestFixture::new();

 let output = Command::new(env!("CARGO_BIN_EXE_flc"))
 .arg(fixture.path())
 .arg("--header")
 .arg("MIT License")
 .output()
 .unwrap();

 assert!(output.status.success());
}

#[test]
fn scan_compliant_files_returns_zero() {
 let fixture = TestFixture::new();
 fixture.create_file("main.rs", "// MIT License\n\nfn main() {}");
 fixture.create_file("lib.rs", "// MIT License\n\npub fn hello() {}");

 let output = Command::new(env!("CARGO_BIN_EXE_flc"))
 .arg(fixture.path())
 .arg("--header")
 .arg("MIT License")
 .output()
 .unwrap();

 assert!(output.status.success());
}

#[test]
fn scan_missing_headers_returns_one() {
 let fixture = TestFixture::new();
 fixture.create_file("main.rs", "fn main() {}"); // No header

 let output = Command::new(env!("CARGO_BIN_EXE_flc"))
 .arg(fixture.path())
 .arg("--header")
```

```
.arg("MIT License")
.output()
.unwrap();

assert!(!output.status.success());
assert_eq!(output.status.code(), Some(1));
}

#[test]
fn scan_respects_gitignore() {
 let fixture = TestFixture::new();
 fixture.create_gitignore(&["ignored/"]);
 fixture.create_file("main.rs", "// MIT License\n\nfn main() {}");
 fixture.create_file("ignored/bad.rs", "fn bad() {}"); // Would fail, but ignored

 let output = Command::new(env!("CARGO_BIN_EXE_flc"))
 .arg(fixture.path())
 .arg("--header")
 .arg("MIT License")
 .output()
 .unwrap();

 assert!(output.status.success());
}

#[test]
fn scan_skips_binary_files() {
 let fixture = TestFixture::new();
 fixture.create_file("main.rs", "// MIT License\n\nfn main() {}");

 // Create binary file
 let binary_path = fixture.dir.path().join("image.png");
 std::fs::write(&binary_path, &[0x89, 0x50, 0x4E, 0x47, 0x00]).unwrap();

 let output = Command::new(env!("CARGO_BIN_EXE_flc"))
 .arg(fixture.path())
 .arg("--header")
 .arg("MIT License")
 .output()
 .unwrap();

 assert!(output.status.success());
}
```

3. tests/integration/fix\_mode.rs:

```
mod common;

use std::fs;
use std::process::Command;
use common::TestFixture;

#[test]
fn fix_adds_header_to_missing_file() {
 let fixture = TestFixture::new();
 let file = fixture.create_file("main.rs", "fn main() {}");

 let output = Command::new(env!("CARGO_BIN_EXE_flc"))
 .arg(fixture.path())
 .arg("--fix")
 .arg("--header")
 .arg("MIT License")
 .output()
 .unwrap();

 assert!(output.status.success());

 let content = fs::read_to_string(&file).unwrap();
 assert!(content.starts_with("// MIT License"));
 assert!(content.contains("fn main() {}"));
}

#[test]
fn fix_preserves_shebang() {
 let fixture = TestFixture::new();
 let file = fixture.create_file("script.sh", "#!/bin/bash\necho hello");

 let output = Command::new(env!("CARGO_BIN_EXE_flc"))
 .arg(fixture.path())
 .arg("--fix")
 .arg("--header")
 .arg("MIT License")
 .output()
 .unwrap();

 assert!(output.status.success());

 let content = fs::read_to_string(&file).unwrap();
 assert!(content.starts_with("#!/bin/bash\n"));
}
```

```

 assert!(content.contains("# MIT License"));
}

#[test]
fn fix_is_idempotent() {
 let fixture = TestFixture::new();
 let file = fixture.create_file("main.rs", "fn main() {}");

 // First fix
 Command::new(env!("CARGO_BIN_EXE_flc"))
 .arg(fixture.path())
 .arg("--fix")
 .arg("--header")
 .arg("MIT License")
 .output()
 .unwrap();

 let content_after_first = fs::read_to_string(&file).unwrap();

 // Second fix
 Command::new(env!("CARGO_BIN_EXE_flc"))
 .arg(fixture.path())
 .arg("--fix")
 .arg("--header")
 .arg("MIT License")
 .output()
 .unwrap();

 let content_after_second = fs::read_to_string(&file).unwrap();

 // Should be identical
 assert_eq!(content_after_first, content_after_second);
}

Step 5.3: Create Performance Benchmark {#step-53-create-performance-benchmark }

Cursor Prompt:
```

Create comprehensive benchmarks in benches/scan\_performance.rs using criterion.

The benchmarks must:

1. Create REALISTIC file content (500 bytes - 5KB, not empty or tiny files)
2. Include proper file distribution:
  - 50% files with correct license headers
  - 40% files missing headers
  - 5% files with shebangs (Python/Shell scripts)
  - 5% binary files (should be skipped)
3. WARM the filesystem cache before timing (read all files once)
4. Test multiple scenarios:
  - scan\_1k\_files (small project baseline)
  - scan\_10k\_files (medium project)
  - scan\_100k\_files (large monorepo - target: < 1 second)
  - single file checks (micro-benchmark)
  - binary detection performance
  - shebang detection performance

\*\*Key Implementation Details:\*\*

```
```rust
// File content distribution for realistic benchmarks
enum FileType {
    WithHeader,      // 50% - has correct license header
    WithoutHeader,   // 40% - missing license header
    Shebang,         // 5% - script with shebang line
    Binary,          // 5% - binary file (should be skipped)
}

// Creates test files with REALISTIC sizes (500-5000 bytes)
fn create_test_files(dir: &TempDir, count: usize) {
    for i in 0..count {
        // Deterministic file size: 500 to 5000 bytes
        let file_size = 500 + ((i * 7919) % 4500);

        // Determine type based on distribution
        let file_type = match i % 20 {
            0..=9 => FileType::WithHeader,      // 50%
            10..=17 => FileType::WithoutHeader, // 40%
            18 => FileType::Shebang,           // 5%
            _ => FileType::Binary,            // 5%
        };

        let content = generate_file_content(file_type, ext, file_size, i);
        fs::write(path, content).ok();
    }
}

/// CRITICAL: Warm the filesystem cache before benchmarking
fn warm_cache(dir: &Path) {
    // Walk all files and read them into memory
    // This ensures consistent "warm cache" measurements
    fn walk_and_read(path: &Path) {
        if path.is_file() {
            let mut file = fs::File::open(path)?;
            let mut buffer = [0u8; 8192];
            while file.read(&mut buffer)? > 0 {}
        } else if path.is_dir() {
            for entry in fs::read_dir(path)?.flatten() {
                walk_and_read(&entry.path());
            }
        }
    }
}
```

```

    }
}

walk_and_read(dir);
}

/// Main benchmark with warm cache
fn bench_scan_sizes(c: &mut Criterion) {
    let mut group = c.benchmark_group("scan");
    group.sample_size(10);

    for count in [1_000, 10_000, 100_000] {
        let dir = TempDir::new().unwrap();
        create_test_files(&dir, count);

        // WARM THE CACHE before benchmarking
        warm_cache(dir.path());

        group.bench_with_input(
            BenchmarkId::new("files_warm", count),
            &count,
            |b, _| {
                b.iter(|| {
                    let scanner = Scanner::new(dir.path(), config).unwrap();
                    scanner.scan()
                });
            }
        );
    }
    group.finish();
}

```

Performance Targets:

Scenario	File Count	Target Time	What It Tests
Small project	1,000	< 50ms	Baseline overhead
Medium project	10,000	< 200ms	Linear scaling
Large monorepo	100,000	< 1,000ms	PRIMARY TARGET
Single file	1	< 0.1ms	Header detection algorithm
Binary detection	1	< 0.01ms	memchr NULL byte search

Why Realistic File Sizes Matter:

File Size	Read Time	Buffer Alloc	What It Tests
Empty (0 bytes)	~0	None	Only walk speed (misleading)
Tiny (~50 bytes)	~0	Minimal	Partial algorithm
Realistic (500-5KB)	Real I/O	Real buffers	Actual performance

Run with: cargo bench

Phase 6: CI/CD Setup (Day 8)

Step 6.1: Create GitHub Actions Workflow

Cursor Prompt:

```
Create .github/workflows/ci.yml:
```

```
```yaml
name: CI

on:
 push:
 branches: [main, develop]
 pull_request:
 branches: [main]

env:
 CARGO_TERM_COLOR: always
 RUST_BACKTRACE: 1

jobs:
 quality-gate:
 name: Quality Gate
 runs-on: ubuntu-latest
 steps:
 - uses: actions/checkout@v4

 - name: Install Rust
 uses: dtolnay/rust-toolchain@stable
 with:
 components: clippy, rustfmt

 - name: Cache cargo
 uses: Swatinem/rust-cache@v2

 - name: Check formatting
 run: cargo fmt --all -- --check

 - name: Clippy (strict)
 run: cargo clippy --workspace --all-targets -- -D warnings

 - name: Install cargo-deny
 run: cargo install cargo-deny

 - name: Security audit
 run: cargo deny check

 - name: Run tests
```

```
run: cargo test --workspace

- name: Run benchmarks (compile only)
 run: cargo bench --no-run

coverage:
 name: Code Coverage
 runs-on: ubuntu-latest
 steps:
 - uses: actions/checkout@v4

 - name: Install Rust
 uses: dtolnay/rust-toolchain@stable
 with:
 components: llvm-tools-preview

 - name: Install cargo-llvm-cov
 run: cargo install cargo-llvm-cov

 - name: Generate coverage
 run: cargo llvm-cov --workspace --lcov --output-path lcov.info

 - name: Upload toCodecov
 uses: codecov/codecov-action@v3
 with:
 files: lcov.info

cross-platform:
 name: Test on ${{ matrix.os }}
 runs-on: ${{ matrix.os }}
 strategy:
 matrix:
 os: [ubuntu-latest, macos-latest, windows-latest]
 steps:
 - uses: actions/checkout@v4

 - name: Install Rust
 uses: dtolnay/rust-toolchain@stable

 - name: Run tests
 run: cargo test --workspace

release:
```

```
name: Release Build
if: startsWith(github.ref, 'refs/tags/')
needs: [quality-gate, cross-platform]
runs-on: ${{ matrix.os }}
strategy:
 matrix:
 include:
 - os: ubuntu-latest
 target: x86_64-unknown-linux-musl
 artifact: flc-linux-x86_64
 - os: macos-latest
 target: x86_64-apple-darwin
 artifact: flc-macos-x86_64
 - os: macos-latest
 target: aarch64-apple-darwin
 artifact: flc-macos-aarch64
 - os: windows-latest
 target: x86_64-pc-windows-msvc
 artifact: flc-windows-x86_64.exe
steps:
 - uses: actions/checkout@v4

 - name: Install Rust
 uses: dtolnay/rust-toolchain@stable
 with:
 targets: ${{ matrix.target }}

 - name: Build release
 run: cargo build --release --target ${{ matrix.target }}

 - name: Upload artifact
 uses: actions/upload-artifact@v3
 with:
 name: ${{ matrix.artifact }}
 path: target/${{ matrix.target }}/release/flc*
```

---

```
Step 6.2: Create README {#step-62-create-readme }
```

\*\*Cursor Prompt:\*\*

Create [README.md](#) with comprehensive documentation:

# Fast License Checker (flc)

Blazing-fast license header verification for your codebase.

## Features

- **Fast:** Scans 100,000 files in under 1 second
- **Git-aware:** Automatically respects .gitignore
- **Safe:** Never corrupts binary files
- **Smart:** Handles shebangs, XML declarations, and various comment styles
- **CI-ready:** JSON and GitHub Actions output formats

## Installation

### From source

```
```bash
cargo install fast-license-checker
````
```

### Pre-built binaries

Download from [Releases](#)

## Usage

### Scan mode (check for missing headers)

```
```bash
flc . # Scan current directory
flc src/ # Scan specific directory
flc --header "MIT License" . # Specify header text
flc --license LICENSE.txt . # Specify header from file
````
```

## Fix mode (add missing headers)

```
```bash
flc --fix . # Add headers to files missing them
```

```

## CI Integration

```
```bash
flc --output github . # GitHub Actions annotations
flc --output json . # JSON for programmatic use
```

```

## Configuration

Create `.license-checker.toml` in your project root:

```
```toml
license_header = """
Copyright (c) 2024 Your Name
Licensed under the MIT License
"""

[comment_styles]
rs = { prefix = "//" }
py = { prefix = "#" }
html = { prefix = "" }

ignore_patterns = ["vendor/", "generated/"]
```

```

## Exit Codes

- 0 : All files have valid headers (scan) or fixes applied (fix)
- 1 : Missing headers found (scan) or errors encountered

# Performance

| Files   | Time (warm cache) |
|---------|-------------------|
| 1,000   | ~10ms             |
| 10,000  | ~100ms            |
| 100,000 | ~800ms            |

# License

MIT OR Apache-2.0

## Cursor Workflow Summary

| Phase        | You (Architect)        | Cursor (Implementation)                  |
|--------------|------------------------|------------------------------------------|
| 1. Bootstrap | Define structure, deps | Generate Cargo.toml, configs             |
| 2. RFC       | Write requirements     | Generate RFC, diagrams                   |
| 3. Design    | Review types, approve  | Generate design_sketch.rs                |
| 4. Core      | Review each module     | Implement types, scanner, checker, fixer |
| 5. CLI       | Define UX, args        | Implement clap, reporters                |
| 6. Tests     | Define test cases      | Generate unit + integration tests        |
| 7. CI/CD     | Define quality gates   | Generate workflows                       |

# Key Architecture Decisions

| Decision                                       | Context                                                | Trade-offs                                    |
|------------------------------------------------|--------------------------------------------------------|-----------------------------------------------|
| Use <code>ignore</code> crate                  | Need <code>.gitignore</code> support, parallel walking | Adds dependency, but battle-tested in ripgrep |
| Use <code>rayon</code> over <code>tokio</code> | File scanning is CPU-bound                             | Can't use async, but simpler model            |
| Use byte slices<br><code>&amp;[u8]</code>      | Binary detection, performance                          | Must handle encoding explicitly               |
| Use <code>thiserror + anyhow</code>            | Typed lib errors, friendly CLI                         | Two error handling styles                     |
| Atomic writes via rename                       | Prevent corruption on crash                            | Requires temp file creation                   |
| Max 8KB header check                           | Don't read entire large files                          | Could miss late headers (rare)                |

## Validation Checkpoints

After each phase, verify:

- [ ] cargo check passes
- [ ] cargo fmt --check passes
- [ ] cargo clippy -- -D warnings passes
- [ ] cargo test passes
- [ ] No unwrap/expect/panic in non-test code
- [ ] All public items have doc comments
- [ ] cargo doc generates without warnings